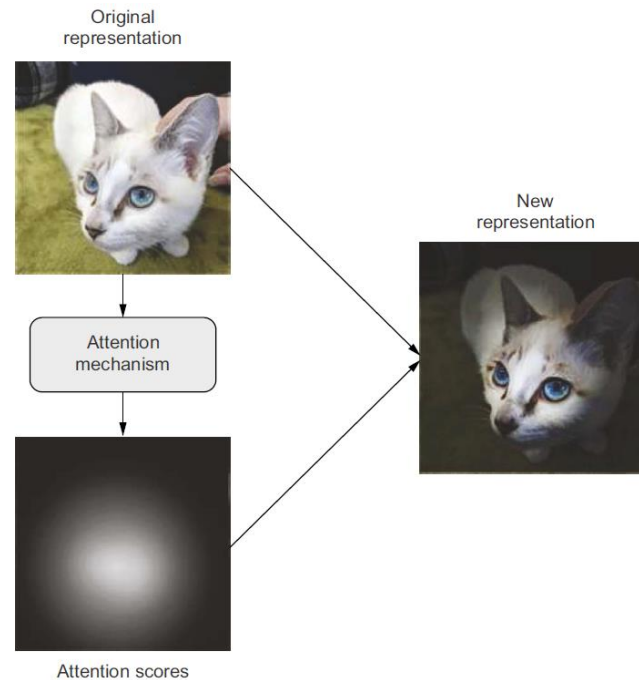


Module 7 – Part 1

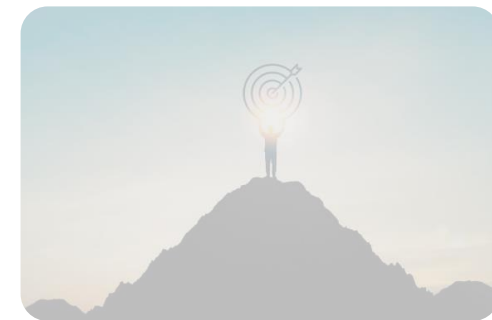
Transformers prerequisites

Why Attention is **ALL** you need?



➔ Road map!

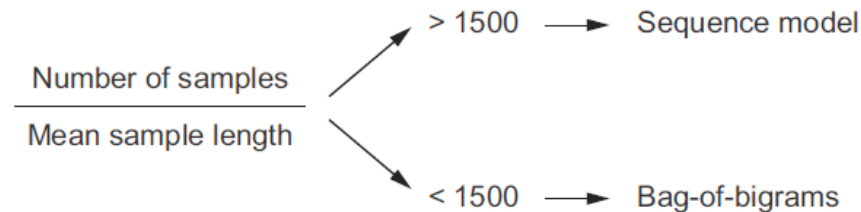
- Module 1- Introduction to Deep Learning
- Module 2- Setting up Deep Learning Environment
- Module 3- Machine Learning review (ML fundamentals + models)
- Module 4- Deep Neural Networks (NN and DNN)
- Module 5- Deep Computer Vision (CNN, R-CNN, YOLO, FCN)
- Module 6- Deep Sequence Modeling (RNN, LSTM)
- **Module 7- Transformers (Attention is all you need!)**
- Module 8- Deep Generative Modeling (AE, VAE, GAN)
- Module 9- Deep Reinforcement Learning (DQN, PG)




→ Transformers for NLP

- Starting in 2017 (**Attention is all you need!**), transformers started overtaking RNN across most NLP tasks.
- NLP architecture depends on word representation method
 - **Discard order** and treat text as an unordered set of words → bag-of-words models
 - **Respect order** and treat words one at a time (steps in timeseries) → recurrent models
- When to use sequence model over bag-of-words?

- For text classification →



- For any other NLP task → Transformers



Transformers vs other sequence models

- Transformer architecture is technically **order-agnostic**, yet it **injects word-position** information into the representations it processes (**hybrid approach**)
- Transformers simultaneously look at different parts of a sentence (unlike RNNs) while still being order-aware.

The cat, sat on the mat.

| NLP Models | Word order awareness | Context awareness (cross-word interactions) |
|-----------------|----------------------|---|
| Bag of unigrams | No | No |
| Bag of Bigrams | Very limited | No |
| RNN | Yes | No |
| Transformer | Yes | Yes |

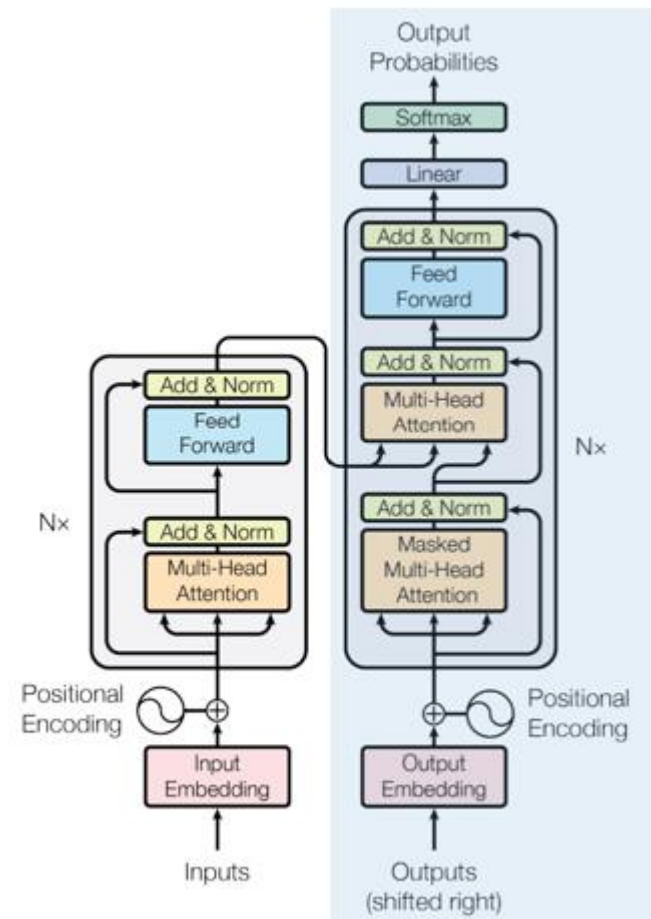
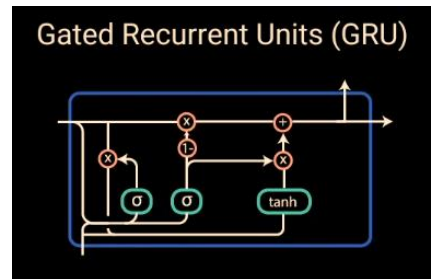
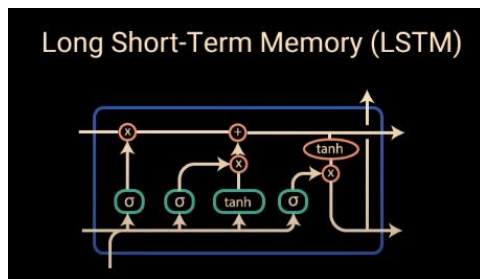
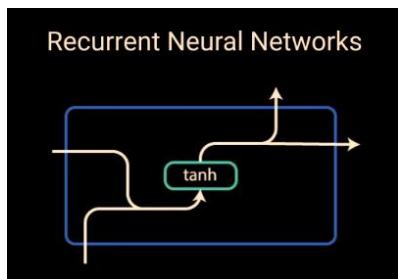




Sequence Modeling Design Criteria

To model sequence data efficiently, we need an architecture that:

- Preserve the **order**
- Account for **long-term dependencies**
- Handle different **input-length**
- **Share parameters** across the sequence



➔ Applications of Transformers


- Transformers are taking NLP, Computer vision and reinforcement learning by storm.
- NLP applications:
 - Machine translation, text generation, text summarization, text classification, chatbots, questions answering etc.
 - BERT, GPT
- Computer vision applications:
 - Image captioning, object detection and segmentation
 - ViT (vision transformers)
- Reinforcement learning applications:
 - Game playing, robotics and autonomous driving





Sequence-to-sequence modeling

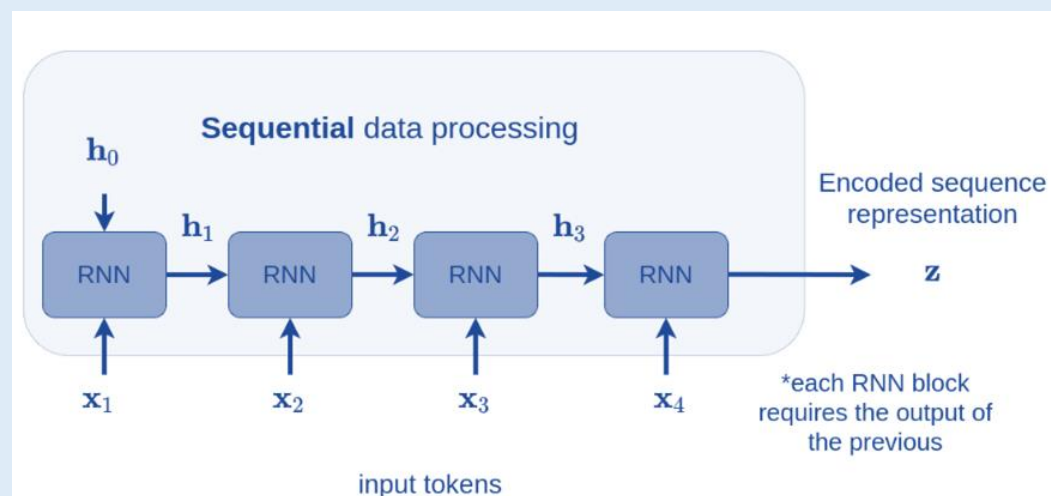
- The goal is to **transform** an input sequence (**source**) to a new one (**target**).




Encoder – Decoder

- Recurrent Neural Networks (RNNs) were the prevailing method for sequence-to-sequence learning until Transformers demonstrated superior performance.

Encoder

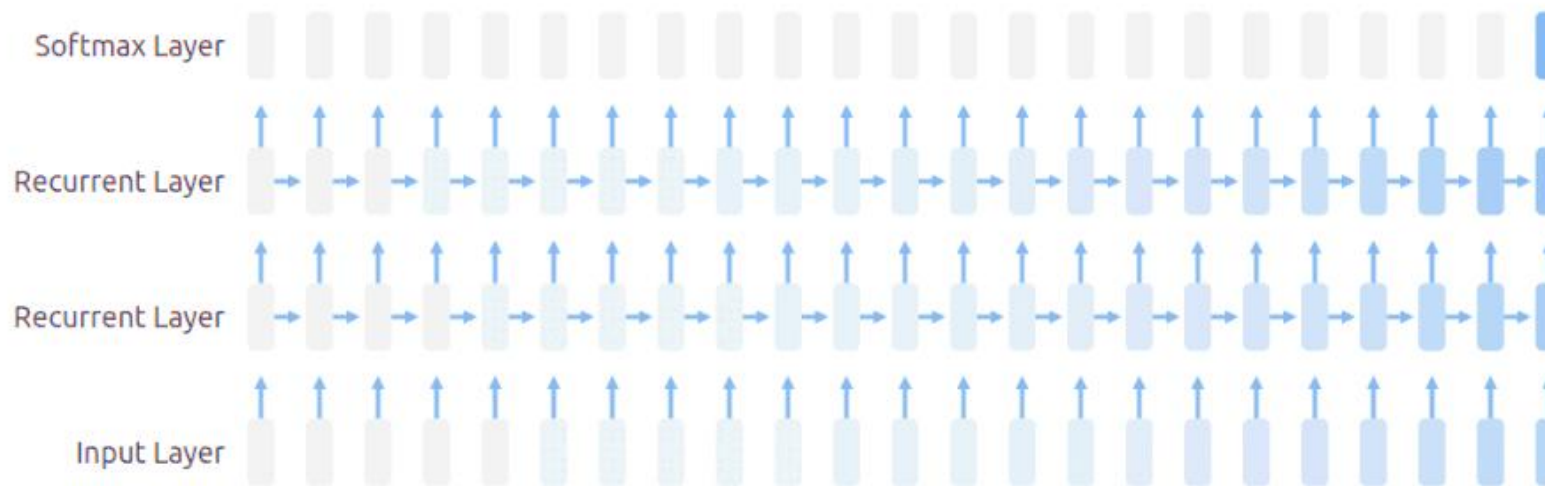


Decoder



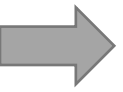
→ Limitations of RNN

- **Bottleneck problem:** the Encoder state vector(s) must store the entire input sequence representation → Significant **limitations on** translatable sentence **size** and **complexity**
- RNN tends to progressively **forget about the past** (~100 tokens) and eventually pays more attention to the **last parts** of the sequence.
- **Vanishing gradient** problem:




Vanishing Gradient: where the contribution from the earlier steps becomes insignificant in the gradient for the vanilla RNN unit. <https://distill.pub/2019/memorization-in-rnns/>

Attention is ALL you need!



Attention

- To avoid **vanishing gradient**, we need to form a **direct connection** with each timestamp.
- By letting the decoder have an attention mechanism, we relieve the encoder from the burden of having to encode all information in the source sentence (**bottleneck problem**)
- Attention is a mechanism that allows the model to **weigh (score)** and focus on specific parts of the input when generating output.
- Attention mechanism can be applied to any encoder and decoder architecture (RNN, LSTM, GRU, CNN, etc)



→ Attention, deep dive

- What does it mean **mathematically**?
- **Vocabulary**: collection of symbols in each sequence
- Converting symbols to numbers: **One-hot encoding**
- Dot product can be used to measure **similarity**
- One-hot vectors can pull out a particular row of a matrix!

The diagram shows a 3x4 matrix A (labeled 'A') and a 4x2 matrix B (labeled 'B') being multiplied to produce a 3x2 result matrix. Matrix A has rows [1, 0, 0, 0], [0, 0, 0, 1], and [0, 0, 1, 0]. Matrix B has rows [.2, .9], [.7, 0], [.8, .3], and [.1, .4]. The result matrix has rows [.2, .9], [.1, .4], and [.8, .3]. This demonstrates how a one-hot vector in A selects a specific row from B.

| | | | |
|---|---|---|---|
| 1 | 0 | 0 | 0 |
| 0 | 0 | 0 | 1 |
| 0 | 0 | 1 | 0 |

| | |
|----|----|
| .2 | .9 |
| .7 | 0 |
| .8 | .3 |
| .1 | .4 |


=

| | |
|----|----|
| .2 | .9 |
| .1 | .4 |
| .8 | .3 |

→ First order sequence model

- Example:

- Show me my **directories** please
- Show me my **files** please
- Show me my **photos** please



- **Vocabulary** size = 7 {directories, files, me, my, photos, please, show}.
- Markov chain transition model
- Matrix form

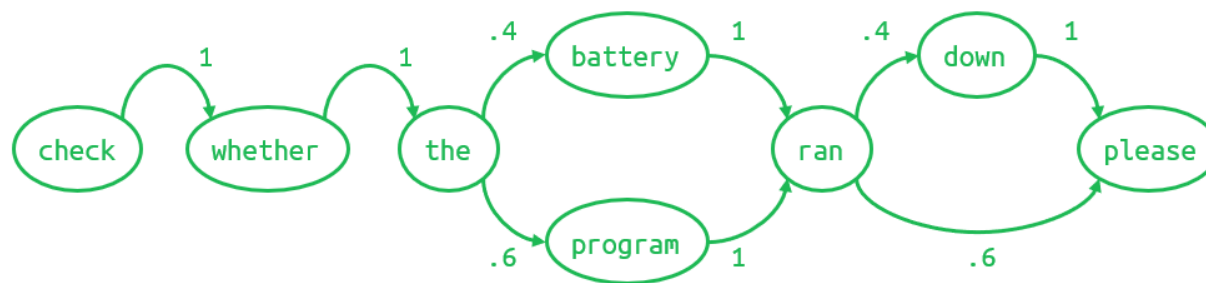
| | directories | files | me | my | photos | please | show |
|-------------|-------------|-------|----|----|--------|--------|------|
| directories | 0 | 0 | 0 | 0 | 0 | 1 | 0 |
| files | 0 | 0 | 0 | 0 | 0 | 1 | 0 |
| me | 0 | 0 | 0 | 1 | 0 | 0 | 0 |
| my | .2 | .3 | 0 | 0 | .5 | 0 | 0 |
| photos | 0 | 0 | 0 | 0 | 0 | 1 | 0 |
| please | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| show | 0 | 0 | 1 | 0 | 0 | 0 | 0 |

=

| directories | files | me | my | photos | please | show |
|-------------|-------|----|----|--------|--------|------|
| .2 | .3 | 0 | 0 | .5 | 0 | 0 |

→ Second order sequence model

- First order Markov model only looks at the single most recent word.
- Predicting based on only one last word is hard! Let's consider the two most recent words!
- Example: (a 40/60 proportion)
 - Check whether the **battery** **ran** **down** please.
 - Check whether the **program** **ran** **please**.
- First order model:
- How can we remove the uncertainty after the word “ran”?





Second order sequence model

- Check whether the **battery ran down** please.
- Check whether the **program ran** please.
- **Vocabulary**: {battery, check, down, please, program, ran, the, whether} **size = 8**

| | battery | check | down | please | program | ran | the | whether |
|---------|---------|-------|------|--------|---------|-----|-----|---------|
| battery | 0 | 0 | 0 | 0 | 0 | 1 | 0 | 0 |
| check | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 1 |
| down | 0 | 0 | 0 | 1 | 0 | 0 | 0 | 0 |
| please | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |
| program | 0 | 0 | 0 | 0 | 0 | 1 | 0 | 0 |
| ran | 0 | 0 | .4 | .6 | 0 | 0 | 0 | 0 |
| the | .4 | 0 | 0 | 0 | .6 | 0 | 0 | 0 |
| whether | 0 | 0 | 0 | 0 | 0 | 0 | 1 | 0 |

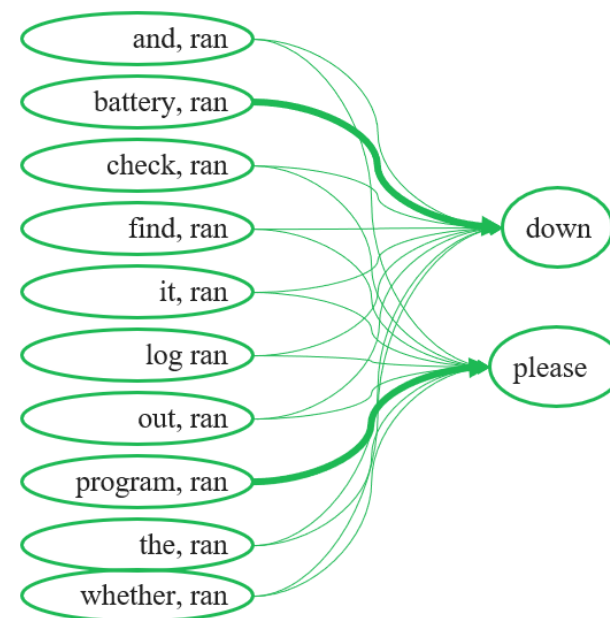
| | battery | check | down | please | program | ran | the | whether |
|---------------|---------|-------|------|--------|---------|-----|-----|---------|
| battery ran | 0 | 0 | 1 | 0 | 0 | 0 | 0 | 0 |
| check whether | 0 | 0 | 0 | 0 | 0 | 0 | 1 | 0 |
| program ran | 0 | 0 | 0 | 1 | 0 | 0 | 0 | 0 |
| the battery | 0 | 0 | 0 | 0 | 0 | 1 | 0 | 0 |
| the program | 0 | 0 | 0 | 0 | 0 | 1 | 0 | 0 |
| ran down | 0 | 0 | 0 | 1 | 0 | 0 | 0 | 0 |
| whether the | .4 | 0 | 0 | 0 | .6 | 0 | 0 | 0 |
| ... | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 |



Higher order sequence models?

- Check the **program** log and find out whether it **ran** please.
- Check the **battery** log and find out whether it **ran** down please.
- What comes after the word “**ran**”? It is unreasonable to investigate 9th order sequence model! (Vocab Size⁹) combinations!
- Solution: Second order sequence model **with skips**

| | and | battery | check | down | find | it | log | out | please | program | ran | the | whether |
|--------------|-----|---------|-------|------|------|----|-----|-----|--------|---------|-----|-----|---------|
| and, ran | | | .5 | | | | | .5 | | | | | |
| battery, ran | | | 1 | | | | | 0 | | | | | |
| check, ran | | | .5 | | | | | .5 | | | | | |
| down, ran | | | 0 | | | | | | | | | | |
| find, ran | | | .5 | | | | | .5 | | | | | |
| it, ran | | | .5 | | | | | .5 | | | | | |
| log, ran | | | .5 | | | | | .5 | | | | | |
| out, ran | | | .5 | | | | | .5 | | | | | |
| please, ran | | | | | | | | | | | | | |
| program, ran | | | 0 | | | | | 1 | | | | | |
| ran, ran | | | | | | | | | | | | | |
| the, ran | | | .5 | | | | | .5 | | | | | |
| whether, ran | | | .5 | | | | | .5 | | | | | |



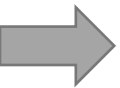


Masking

- **Masking**: Crossing out all the uninformative feature votes
- The only important rows are *battery, ran* and *program, ran*. We could **mask** everything else!

Check the *program* log and find out whether it *ran* please.

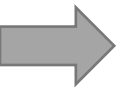
| | and, ran | battery, ran | check, ran | down, ran | find, ran | it, ran | log, ran | out, ran | please, ran | program, ran | ran, ran | the, ran | whether, ran |
|---------------------------|----------|--------------|------------|-----------|-----------|---------|----------|----------|-------------|--------------|----------|----------|--------------|
| feature activities | 1 | 0 | 1 | 0 | 1 | 1 | 1 | 1 | 0 | 1 | 0 | 1 | 1 |
| | * | | | | | | | | | | | | |
| mask | 0 | 1 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 1 | 0 | 0 | 0 |
| | = | | | | | | | | | | | | |
| masked feature activities | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 0 | 1 | 0 | 0 | 0 |



Masking: selective second order model with skips

- The mask has the effect of **hiding** a lot of the transition matrix
- In this example, it hides the combination of **ran** with everything except **battery** and **program**, leaving **just the features that matter**
- This process of selective masking is the **attention** thing!


| | and | battery | check | down | find | it | log | out | please | program | ran | the | whether |
|--------------|-----|---------|-------|------|------|----|-----|-----|--------|---------|-----|-----|---------|
| and, ran | | | | | | | | | | | | | |
| battery, ran | | | 1 | | | | | 0 | | | | | |
| check, ran | | | | | | | | | | | | | |
| down, ran | | | | | | | | | | | | | |
| find, ran | | | | | | | | | | | | | |
| it, ran | | | | | | | | | | | | | |
| log, ran | | | | | | | | | | | | | |
| out, ran | | | | | | | | | | | | | |
| please, ran | | | | | | | | | | | | | |
| program, ran | | | 0 | | | | | 1 | | | | | |
| ran, ran | | | | | | | | | | | | | |
| the, ran | | | | | | | | | | | | | |
| whether, ran | | | | | | | | | | | | | |



Attention as Matrix Multiplication

- Stack the mask vectors for every word into a matrix (**Keys**)
- Use one-hot representation of the most recent word (**Query**) to pull out the relevant mask (**attention score**)
- Wight the tokens in the input sequence (**Values**) by the attention scores to create the **context vector**.

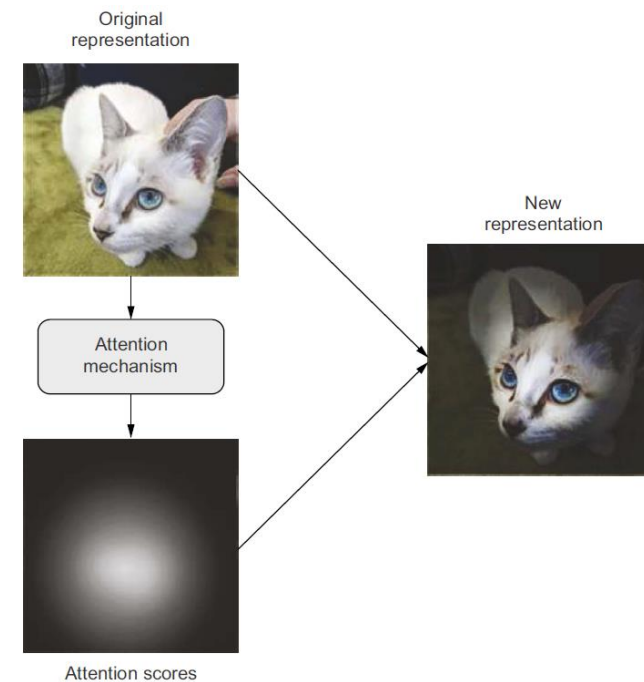
Check the program log and find out whether it *ran* please.

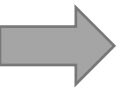




Self-Attention

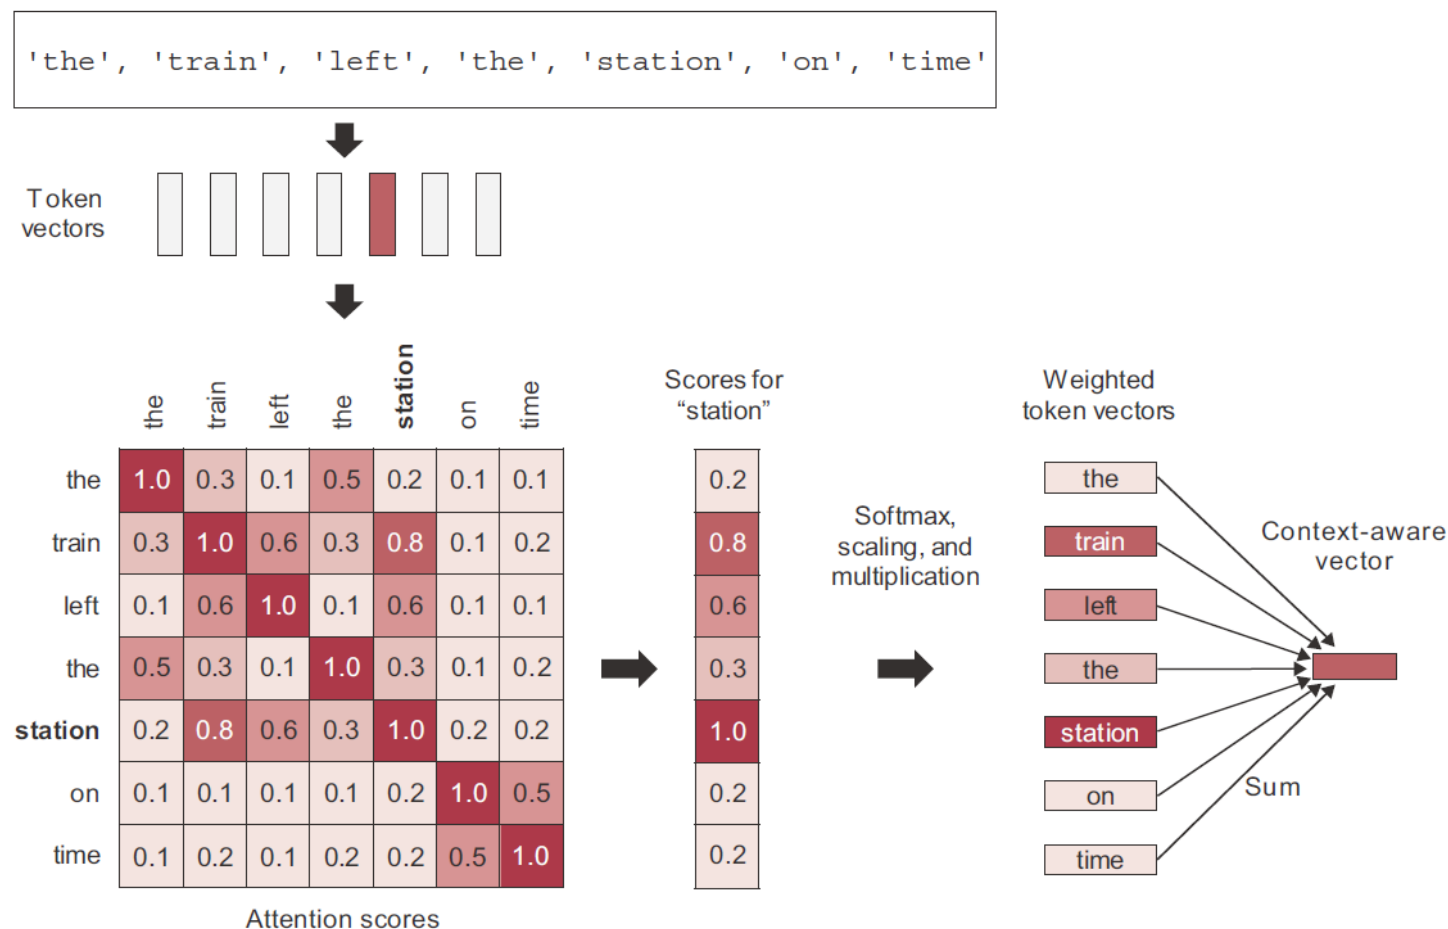
- **Self attention** is a variation of the attention mechanism where the **input and output sequences are the same**, meaning the model is attending to its own input.
- Self-Attention starts by computing **scores** for a set of features. High score → more relevant
- Self attention allows the model to **relate** different parts of the input sequence to each other, capturing dependencies and relationships within the sequence itself
- Have we seen this idea before? Max Pooling and TF-IDF





Self-Attention (context-aware representation)

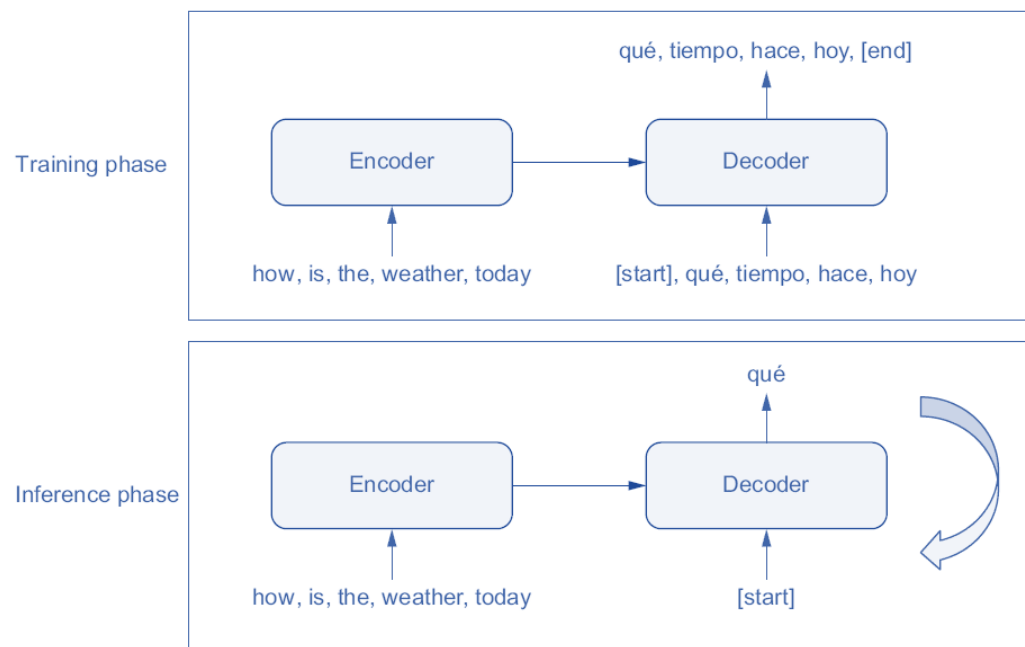
- Self-attention helps to **adjust the representation of a token** by considering the information from **related tokens** in the input sequence → **context awareness**

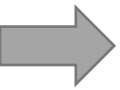




Encoder–Decoder (Machine Translation)

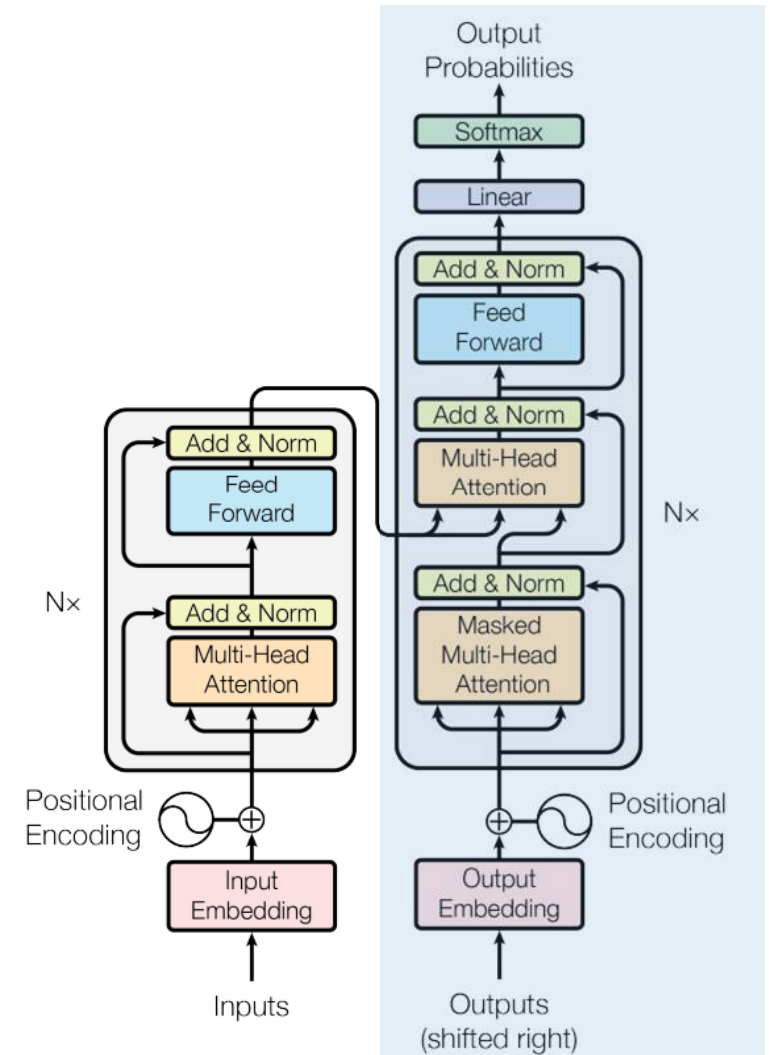
- An **encoder model** turns the source sequence into an **intermediate representation**.
- A **decoder model** is trained to predict the next token in the target sequence by looking at both previous tokens and the encoded source sequence.
- During **inference**, we don't have access to the target sequence (predict it from scratch) → must generate it one token at a time





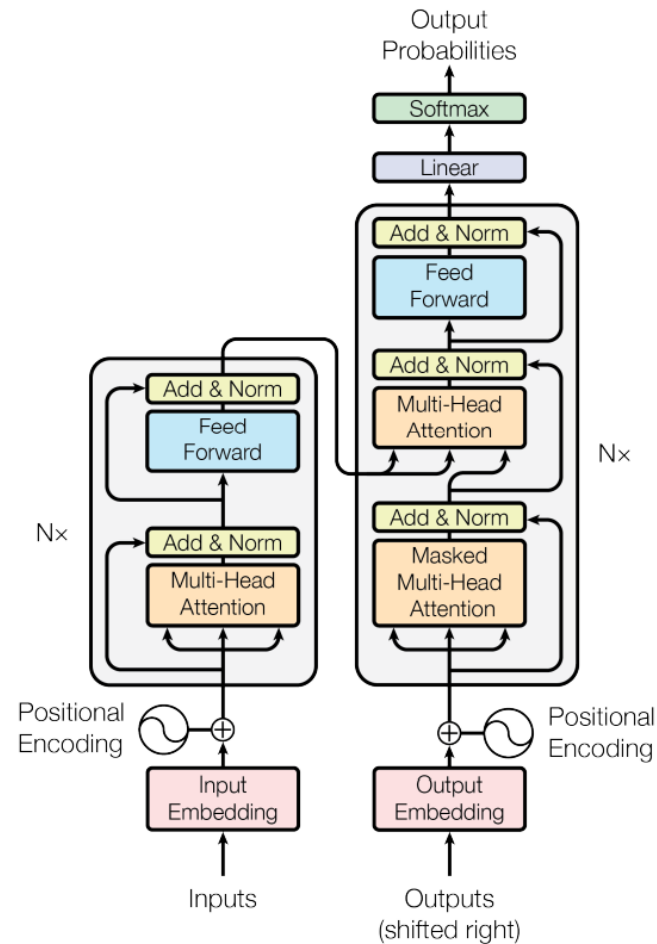
Self-Attention and Transformers

- Transformers **primarily** use **self-attention** mechanisms.
- Both the **encoder** and **decoder** layers apply self-attention to capture relationships and dependencies within the input sequence itself.
- By using self-attention, transformers can:
 - Process inputs in parallel
 - Identify long-range dependencies
 - Model complex relationships more efficiently compared to RNN and CNN.



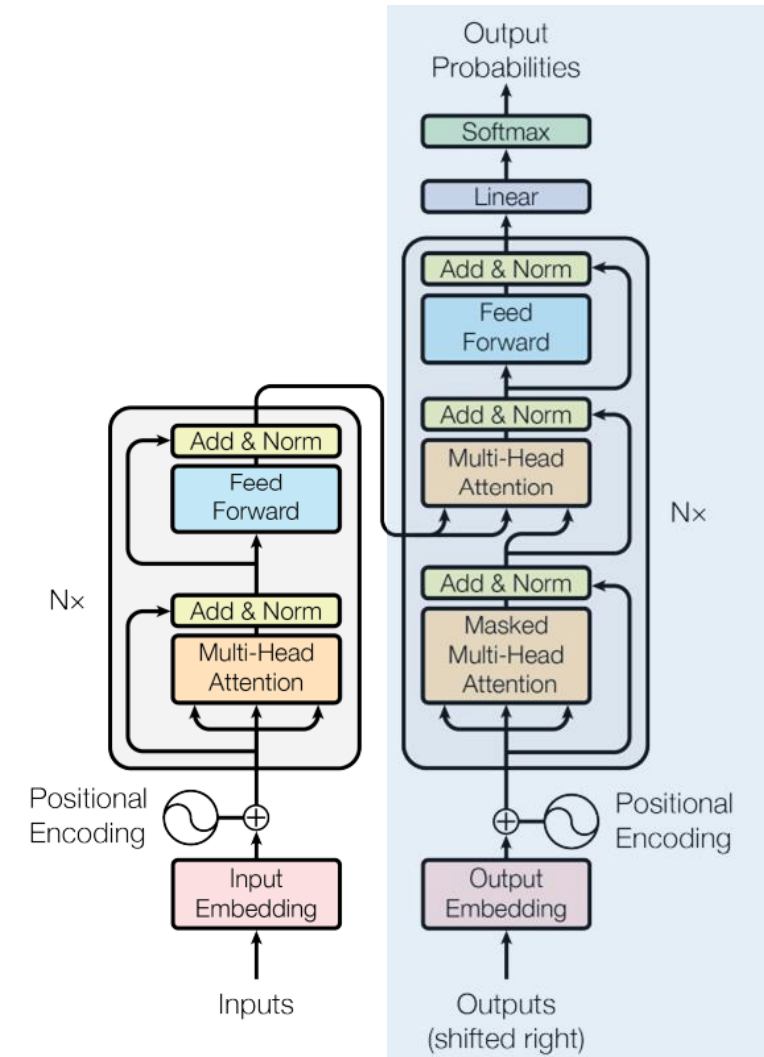
Module 7 – Part 2

Transformers Architecture



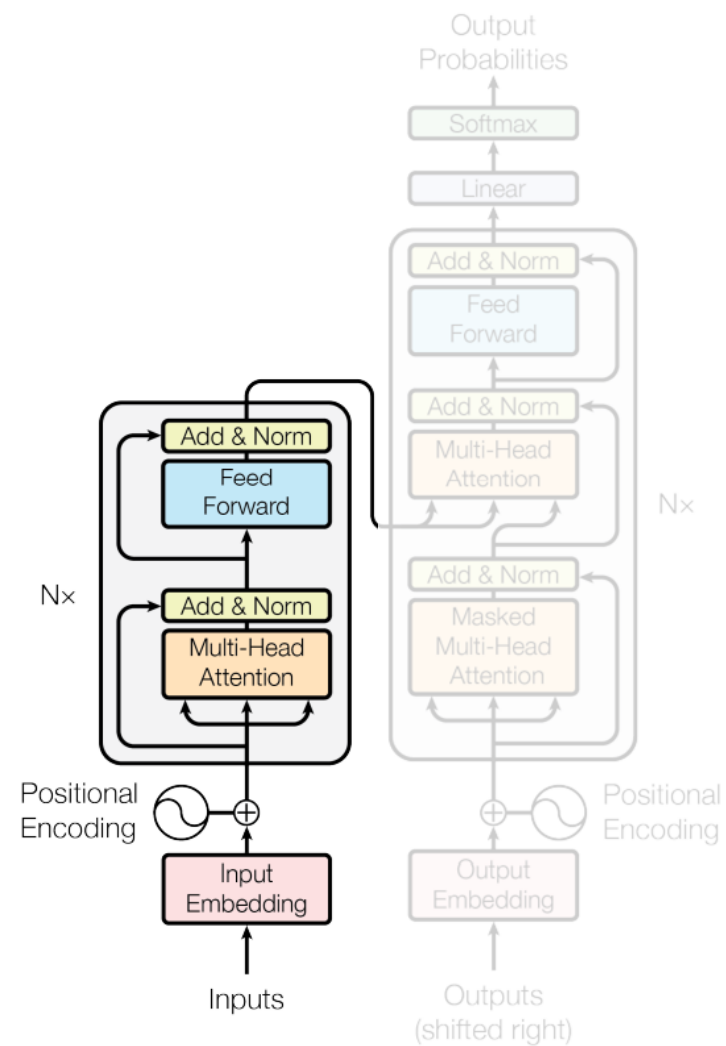
Transformers outline

- Encoder – Decoder architecture
 - Embedding and Positional Encoding
 - Self-Attention (scaled dot product attention)
 - Query-Key-Value model
 - (Masked) Multi-head attention
 - Encoder-Decoder attention
 - Residual connections
 - Layer normalization
 - Feed Forward
 - Softmax layer



Transformer architecture


Encoder






Input Embedding

- **Word embedding** is a numerical representation of words as vectors in a continuous, relatively low-dimensional space compared to the size of the vocabulary, which captures their semantic meaning and **relationships with other words**.
- Word embedding is **learned** through a neural network
- Notion of order is **lost**!
- Example: “ I love Transformers!” with embedding dim=10
 - "I" : [-0.2, 0.1, 0.3, -0.4, 0.5, -0.6, -0.1, -0.3, -0.2, 0.4]
 - "love" : [0.3, -0.2, 0.1, 0.5, -0.4, 0.2, -0.6, -0.1, -0.3, 0.2]
 - "transformers" : [-0.4, -0.3, 0.6, -0.1, 0.2, 0.1, 0.4, 0.5, 0.2, 0.3]
 - "!" : [0.1, -0.4, -0.2, 0.3, 0.2, -0.1, 0.5, -0.3, -0.5, -0.4]



→ Positional Encoding

- Positional encoding **reinject** order information
- positional encoding is a set of small constants, which are added to the word embedding vector before the first self-attention layer.
- If the same word appears in a different position, the actual representation will be slightly different, depending on where it appears in the input sentence
- The model **will figure out** how to leverage this additional information.
- Naïve solution: My name is Pedram $\rightarrow \{0,1,2,3\}$
- Better solution: Add a circular wiggle



Word embedding + positional encoding =
positional embedding



Positional Encoding

- Better solution: Sine – Cosine function

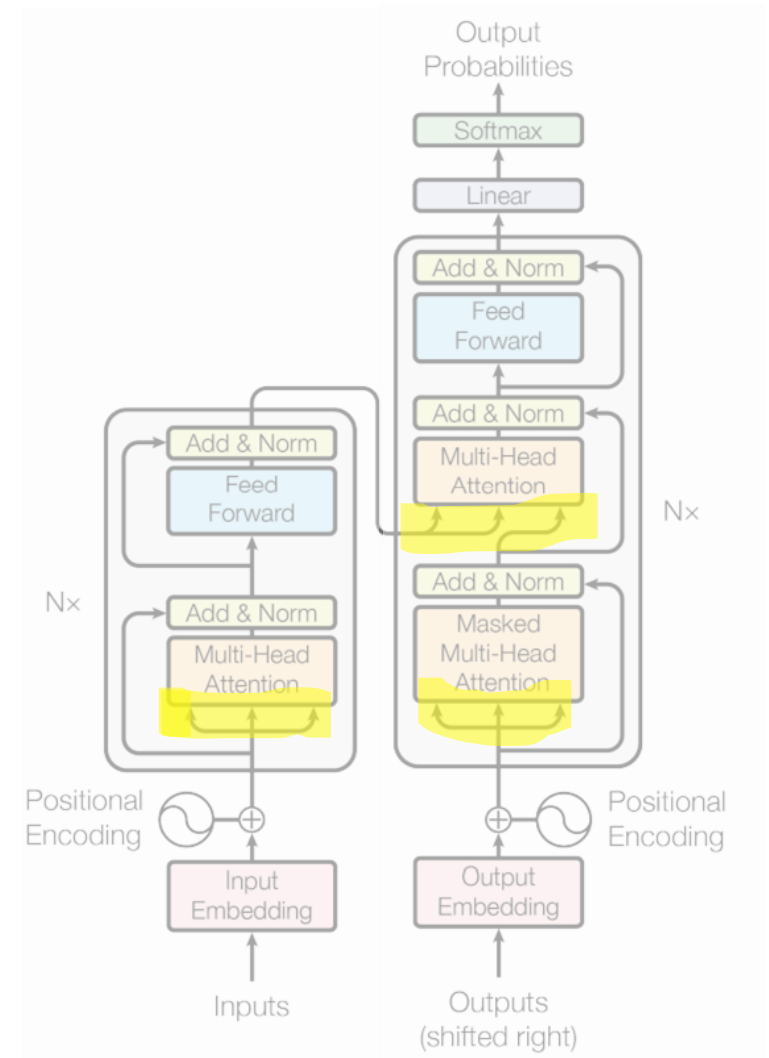
$$PE_{(pos,2i)} = \sin(pos/10000^{2i/d_{\text{model}}})$$

$$PE_{(pos,2i+1)} = \cos(pos/10000^{2i/d_{\text{model}}})$$

$$\vec{p}_t = \begin{bmatrix} \sin(\omega_1 \cdot t) \\ \cos(\omega_1 \cdot t) \\ \sin(\omega_2 \cdot t) \\ \cos(\omega_2 \cdot t) \\ \vdots \\ \sin(\omega_{d/2} \cdot t) \\ \cos(\omega_{d/2} \cdot t) \end{bmatrix}_{d \times 1}$$




Self-Attention mechanism Query-Key-Value



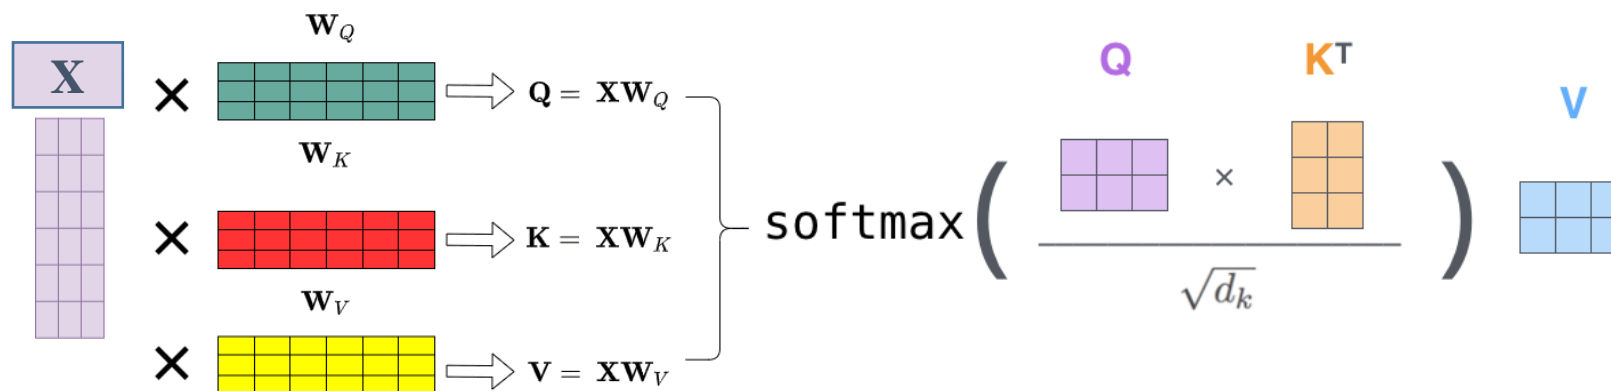

→ The Query-Key-Value

- This terminology comes from **search engines**
- **Query**: What we are looking for?
- **Value**: Body of knowledge that we are trying to extract information from (database)
- **Key**: Set of “keywords” that describes the value in a format that can be readily compared to a query.
- The “query” is **compared** to a set of “keys,” and the match scores are used to **rank** “values” (images in this example).



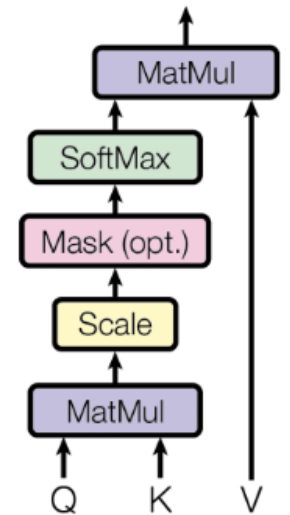
→ Scaled dot product attention

- **Goal:** Identify which part to attend to and Extracting features with high attention
- Scaled dot product attention utilizes three weight matrices, referred to as W_Q , W_K , and W_V which are **learned** as model parameters during training.
- These matrices serve to project the inputs into query, key, and value components of the sequence.



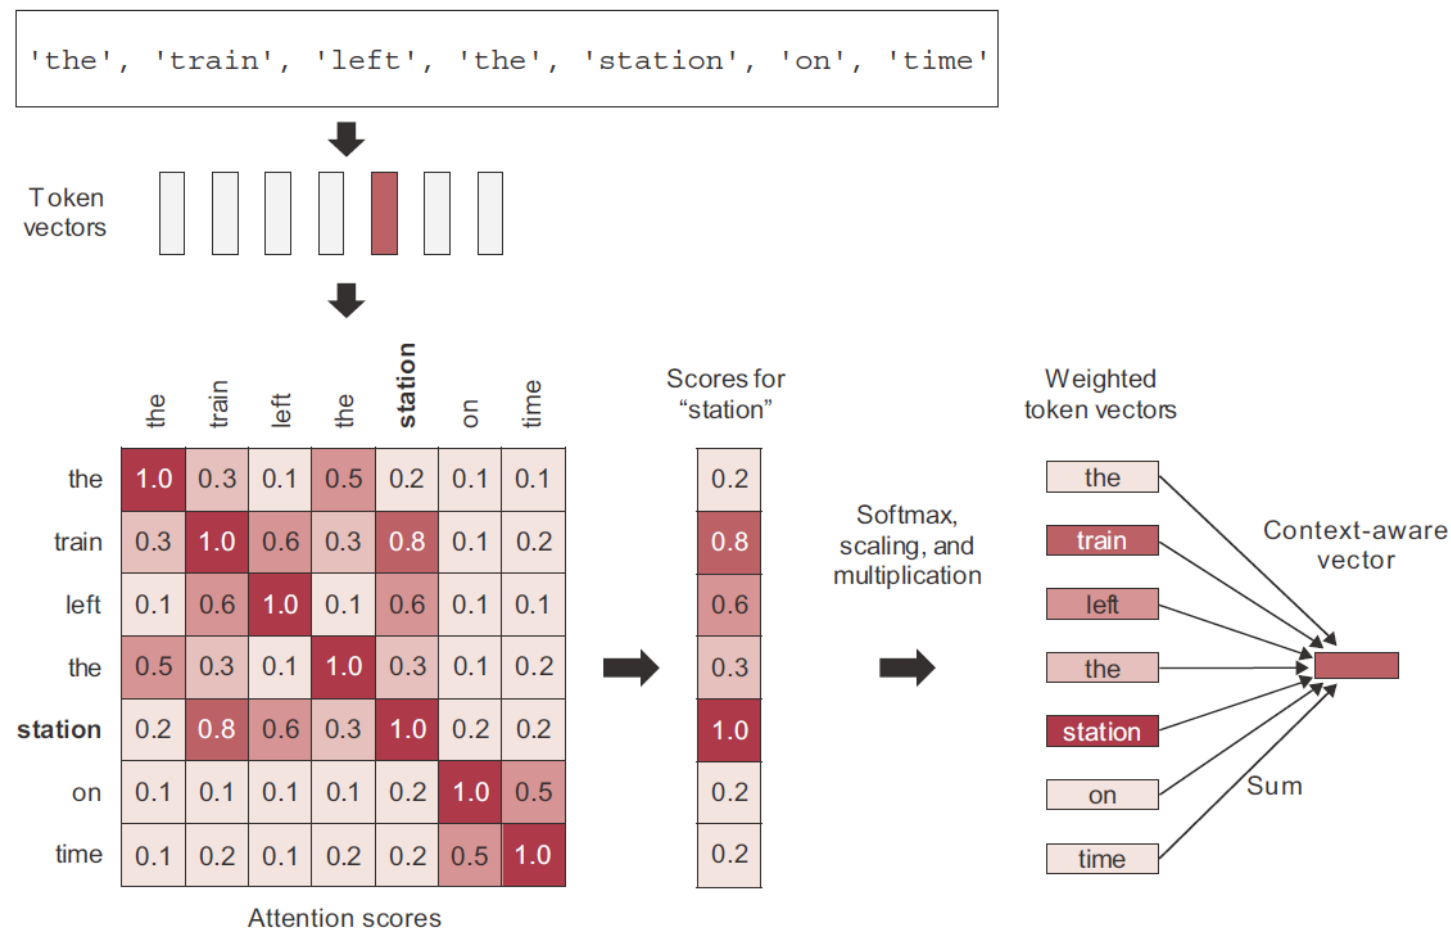
Self-Attention in Language Models

1. The input embeddings are transformed into **query**, **key**, and **value** vectors using different **learned** weight matrices. Each of these vectors is used for different purposes:
 - **Query (Q)**: Represents the current token that the model is focusing on.
 - **Key (K)**: Represents all the tokens in the input sequence, which are compared to the query to compute **attention scores**.
 - **Value (V)**: Represents all the tokens in the input sequence, which are weighted by the attention scores to create the **context vector**.
2. **Attention scores** = $\text{softmax}\left(\frac{Q \cdot K^T}{\sqrt{d_k}}\right)$
3. The attention scores are then used to weight the value vectors, and the weighted sum of these value vectors forms the **context vector** for the current word.
4. This context vector is then used to generate the output.



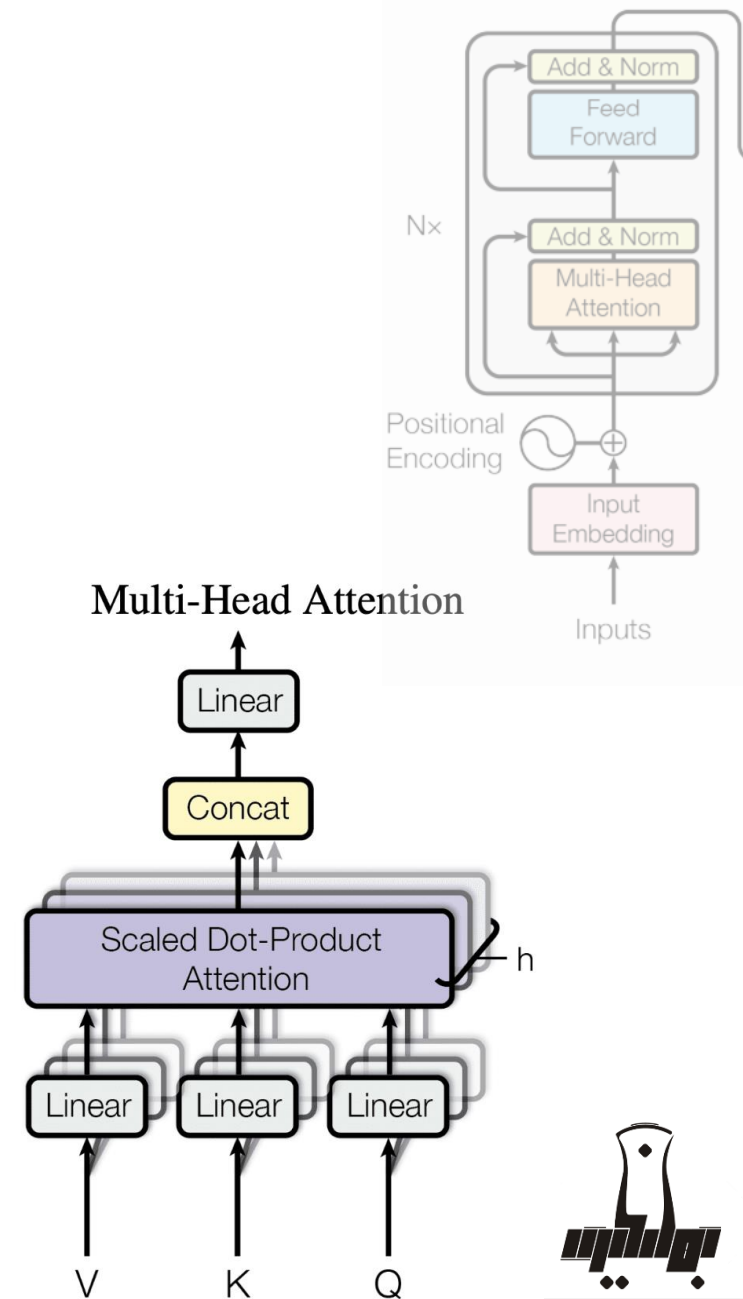


Self-Attention in Language Models



Multi-Head Attention

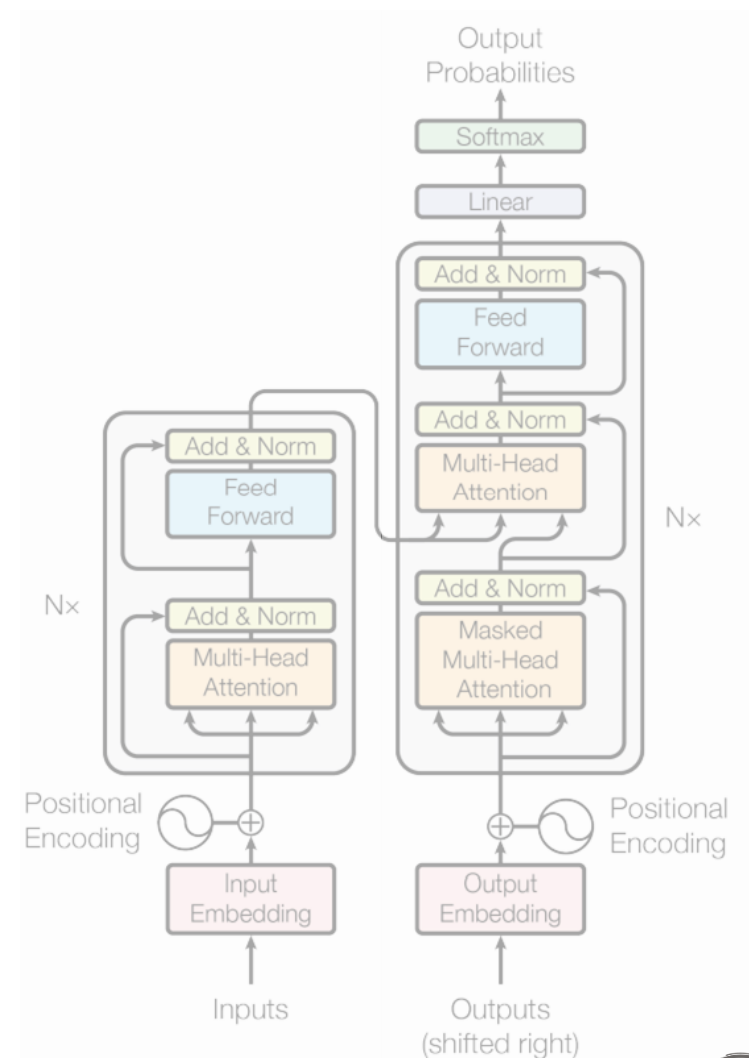
- **Intuition**: it allows the model to **attend to different parts** of the sequence **differently** each time
- Features **within** a self-attention head are **correlated** but mostly independent from those in other heads, similar to Depthwise separable convolutions treating each channel independently to **obtain more expressive representations**.
- Multi-head attention is simply the application of the same idea to self-attention.
- Independent heads help the layer learn different groups of features for each token.

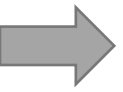




Add & Norm

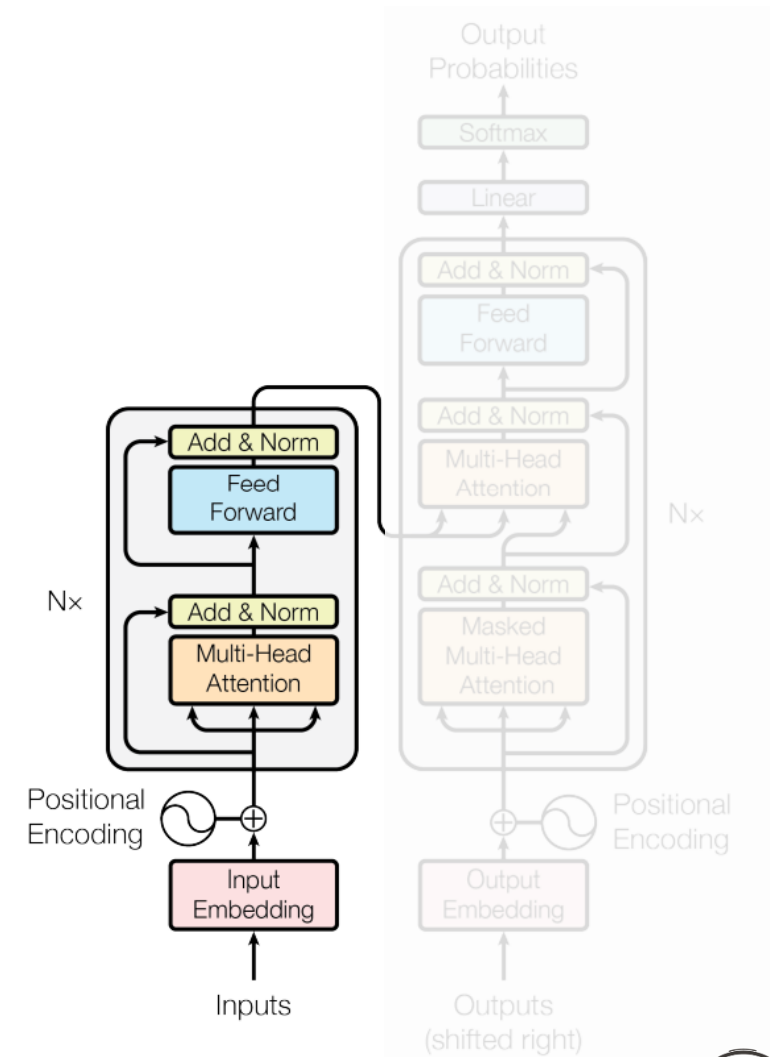
- Residual connections are added to preserve valuable information and prevent information loss during the training process.
- Normalization layers aid gradient flow during backpropagation to improve the training process.
 - Layer Normalization instead of Batch Normalization
 - Normalizing each sequence independently
- Leveraging **standard architectural** patterns such as factoring outputs into multiple independent spaces, adding residual connections, and incorporating normalization layers can enhance the performance of complex models.






Encoder summary

- Input embedding
- Positional encoding
- Encoder block:
 - Multi-head self-attention layer
 - Layer normalization
 - Residual connection
 - MLP (2 linear layers + RELU activation)
 - Second Layer normalization
 - Second residual connection
- Replicating N times (N=6 in the original paper)



Transformer architecture

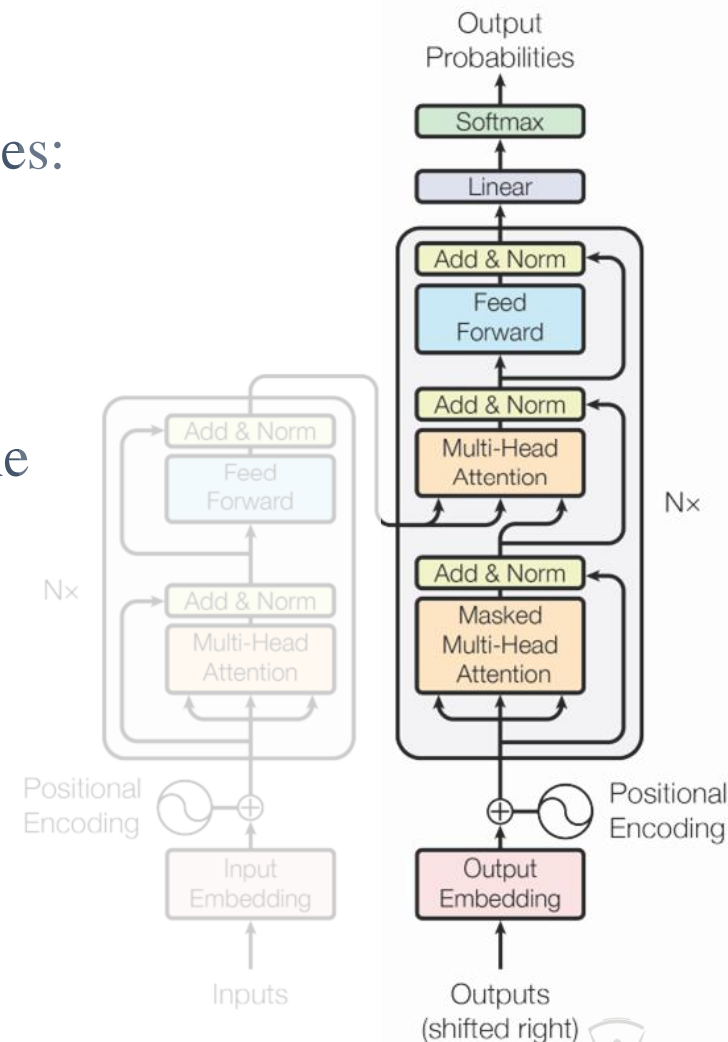
Decoder

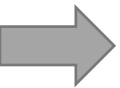




Decoder

- Decoder consist all the components of encoder plus two novel ones:
 - **Masked** multi-head self-attention layer
 - **Encoder-Decoder** Attention layer
- Final decoder block output is passed through a **linear layer** and the probabilities are calculated with a standard **softmax** function
- Decoder is **autoregressive**:
 - Generate output one at a time
 - This is repeated until the **<end>** token is seen.





Masked multi-head self-attention layer

- First Multi-headed attention computes the attention scores for the decoders input.
- For this Multi-headed attention layer, we need to apply **masking** to avoid cheating.

$$\text{MaskedAttention}(\mathbf{Q}, \mathbf{K}, \mathbf{V}) = \text{softmax} \left(\frac{\mathbf{QK}^T + \mathbf{M}}{\sqrt{d_k}} \right) \mathbf{V}$$

| | | | |
|-----|-----|-----|-----|
| 0.7 | 0.1 | 0.1 | 0.1 |
| 0.1 | 0.6 | 0.2 | 0.1 |
| 0.1 | 0.3 | 0.6 | 0.1 |
| 0.1 | 0.3 | 0.3 | 0.3 |

+

| | | | |
|---|------|------|------|
| 0 | -inf | -inf | -inf |
| 0 | 0 | -inf | -inf |
| 0 | 0 | 0 | -inf |
| 0 | 0 | 0 | 0 |

=

| | | | |
|-----|------|------|------|
| 0.7 | -inf | -inf | -inf |
| 0.1 | 0.6 | -inf | -inf |
| 0.1 | 0.3 | 0.6 | -inf |
| 0.1 | 0.3 | 0.3 | 0.3 |

⇒


Softmax(

| | | | |
|-----|------|------|------|
| 0.7 | -inf | -inf | -inf |
| 0.1 | 0.6 | -inf | -inf |
| 0.1 | 0.3 | 0.6 | -inf |
| 0.1 | 0.3 | 0.3 | 0.3 |

)

=


| | | | |
|------|------|------|------|
| 1 | 0 | 0 | 0 |
| 0.37 | 0.62 | 0 | 0 |
| 0.26 | 0.31 | 0.43 | 0 |
| 0.21 | 0.26 | 0.26 | 0.26 |





Encoder-Decoder Attention Layer

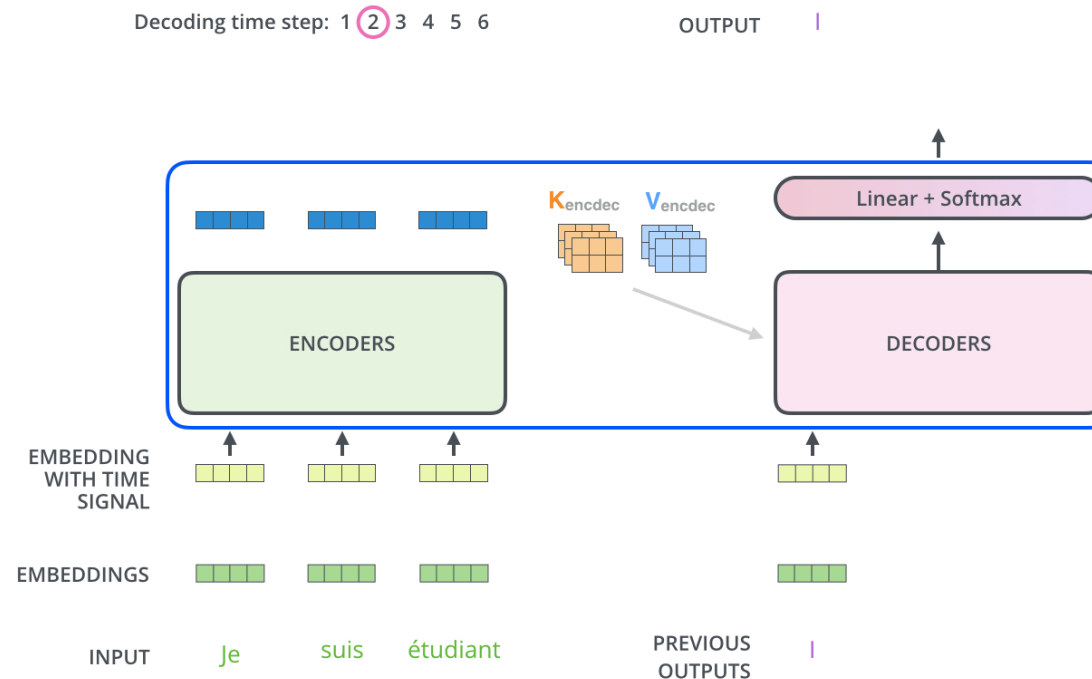
- Transformers also use a form of attention known as "encoder-decoder attention" in the **decoder** layers.
- This is **where the magic happens**, where the decoder processes the encoded representation (K,V)
- This attention mechanism allows the decoder to focus on specific parts of the **input sequence encoded by the encoder** when generating each output element.
- The encoder-decoder attention mechanism helps the model to **align the input and output sequences better**, which is particularly useful in tasks like machine translation.





Encoder-Decoder Attention Layer (cross-attention)

- The encoder output is utilized to generate the **Key** and **Value** matrices.
- On the other hand, the Masked Multi-head attention block's output contains the newly generated sentence, represented as the **Query** matrix in the attention layer.
- This process matched the encoders input to the decoders input allowing the decoder to decide which encoder input is relevant to focus on.



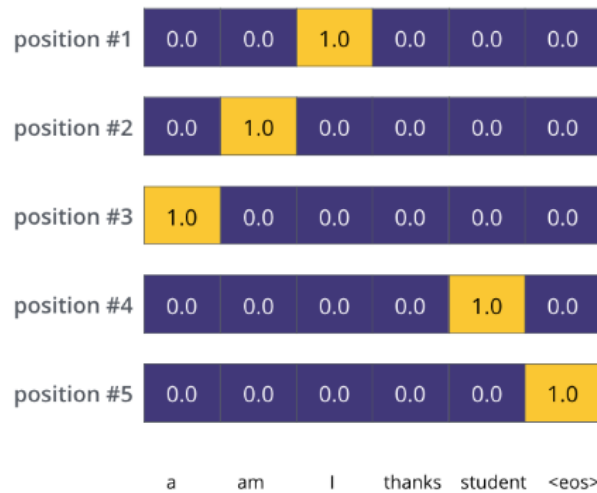


Training

- Like any other deep learning model, training involves:
 - Feed forward, loss computation, backpropagation, parameter updates
- Cross-entropy compare two probability distributions.


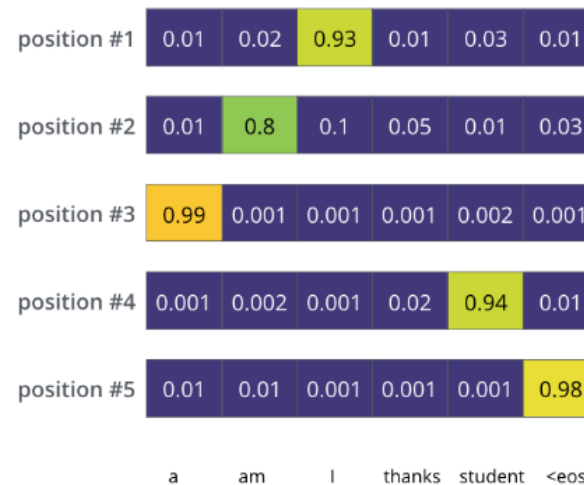
Target Model Outputs

Output Vocabulary: a am I thanks student <eos>



Trained Model Outputs


Output Vocabulary: a am I thanks student <eos>






Decoder summary

- Input embedding
- Positional encoding
- Decoder block:
 - Masked Multi-head self-attention layer
 - First Normalization and residual connection
 - Encoder-Decoder Multi-head attention
 - Second Normalization and residual connection
 - MLP (2 linear layers + RELU activation)
 - Third Normalization and residual connection
- Linear layer followed by a softmax function
- Replicating N times (N=6 in the original paper)



➔ Why transformers work?

- Multi-head attention (multiple representations of the same input)
- Context awareness through self-attention (data-dependent dynamic weights)
- Stacking multiple encoder and decoder blocks (transformer blocks are shape-invariant)





References

- Attention is all you need! *Vaswani et al*
- Deep learning with Python, *Francois Chollet*
- How Transformers work in deep learning and NLP: an intuitive introduction, *Nikolas Adaloglou*
- Transformers from scratch, *Brandon Rohrer*
- The illustrated transformers, *Jay Alammar*

➔ Road map!

- ✓ Module 1- Introduction to Deep Learning
- ✓ Module 2- Setting up Deep Learning Environment
- ✓ Module 3- Machine Learning review (ML fundamentals + models)
- ✓ Module 4- Deep Neural Networks (NN and DNN)
- ✓ Module 5- Deep Computer Vision (CNN, R-CNN, YOLO, FCN)
- ✓ Module 6- Deep Sequence Modeling (RNN, LSTM)
- ✓ Module 7- Transformers (Attention is all you need!)
- Module 8- Deep Generative Modeling (AE, VAE, GAN)
- Module 9- Deep Reinforcement Learning (DQN, PG)

